



Inspiring Excellence

## CSE422: Artificial Intelligence

**Spring 2024**

**Title:** Impact of Player Injuries on Team Performance

**Student Information:**

Name: MD Adnan Hossain Sadi

ID: 22299314

Section: 03

Name: Sarahat Noor Sakib

ID: 19101454

Section: 03

**Submitted to:**

MD. Asif Haider & Mahjabin Chowdhury

Department of Computer Science & Engineering (CSE)

## **Table of contents:**

<b>Content</b>	<b>Page no.</b>
1. Introduction	02
2. Dataset description	03
3. Dataset pre-processing	06
4. Feature scaling	11
5. Dataset splitting	11
6. Model training & testing	13
7. Model selection/Comparison analysis	17
8. Conclusion	22

## **1. Introduction:**

This project aims to analyze the performance of football teams after the return of a key player from injury. In competitive football, the absence of critical players due to injuries can significantly impact team dynamics, strategies, and overall performance. When these players return, their reintegration into the team can either improve or disrupt team cohesion. The project investigates how a team's performance changes in the matches immediately following a player's recovery.

Football teams often rely heavily on key players, and their absence due to injury can affect results. While it's generally assumed that a returning player strengthens the team, this may not always be the case. Factors such as fitness levels, match sharpness, team chemistry, and tactical adjustments can all influence outcomes. The challenge lies in quantifying and understanding these effects to determine whether teams experience performance improvement, stagnation, or decline after a player's return.

The motivation for this project stems from the critical role injuries play in football. Coaches, analysts, and fans often debate the impact of returning players, but these discussions are rarely backed by systematic analysis. By providing data-driven insights, this project aims to support teams in decision-making, such as the timing of a player's return or adjustments to tactics. Additionally, it seeks to deepen our understanding of how individual contributions shape team performance, offering valuable perspectives to players, coaches, and the broader football community.

## 2. Dataset description:

- Source: Kaggle

Link-

<https://www.kaggle.com/datasets/amritbiswas007/player-injuries-and-team-performance-dataset/data>

- **Dataset Description:**

There are in total 42 features in our dataset.

```
1 print(len(data.columns))  
✓ [4] < 10 ms  
42
```

The target feature of our dataset is “Match1\_after\_injury\_Result” which has only three unique values (lose, win & draw). So, the data of the target feature is categorical and the problem we are trying to solve is a classification problem.

There are 656 data points in the dataset.

```
1 print(len(data))  
✓ [5] < 10 ms  
656
```

Our dataset has both quantitative and categorical features.

Below is a breakdown of potential features for this project:

Quantitative Features:

These are numerical values that can be measured or calculated.

Examples include: 'Age', 'FIFA rating', 'Match1\_before\_injury\_GD' etc.

Categorical Features:

These are non-numerical variables that represent distinct categories or groups.

Examples include: 'Position', 'Match1\_before\_injury\_Result' etc.

The following heatmap shows the correlation between various features of our dataset. Diagonal elements are all 1, as each variable is perfectly correlated with itself. Age vs. FIFA rating: Correlation of 0.31, indicating a weak positive relationship. FIFA rating vs. Match1\_before\_injury\_GD: Correlation of 0.15, suggesting a weak positive relationship. Match1\_missed\_match\_GD vs. other variables: Mostly close to 0, indicating very weak or no correlation. Match1\_after\_injury\_GD vs. Match1\_before\_injury\_GD: A correlation of 0.03, indicating negligible relationship. We can see that most of the features of our dataset are not strongly correlated. The only obvious correlation we can observe is between 'Age' and 'FIFA rating' which will not help us that much as none of these are the target feature. This lack of correlation might result in poor accuracy when we train our machine learning model with these dataset.

```

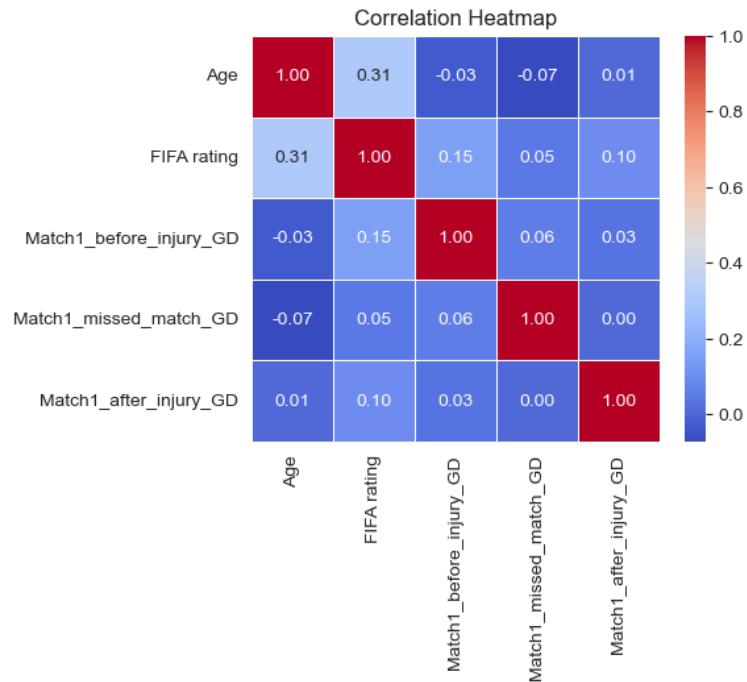
1 input_features = ['Age', 'FIFA rating', 'Match1_before_injury_GD', 'Match1_missed_match_GD']
2 output_features = ['Match1_after_injury_GD']
3 selected_features = input_features + output_features
4 correlation_matrix = data[selected_features].corr()
✓ [51] < 10 ms

```

```

1 plt.figure(figsize=(5, 4))
2 sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
3 plt.title('Correlation Heatmap')
4 plt.show()
✓ [52] 119ms

```

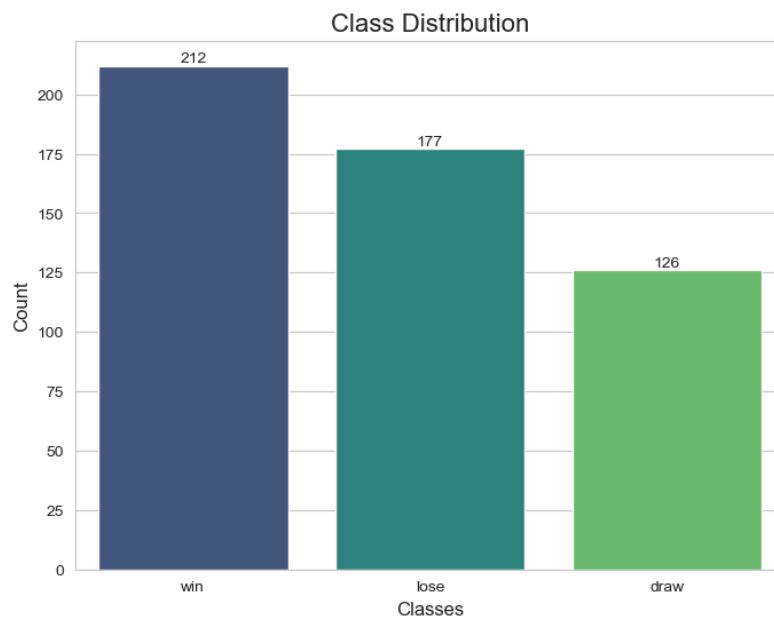


- Imbalance Dataset:** The dataset we used for this project is an imbalance dataset as the target feature does not have equal instances of all the unique classes.

```

1 target_column = 'Match1_after_injury_Result'
2
3 class_counts = data[target_column].value_counts()
4
5 plt.figure(figsize=(8, 6))
6 sns.barplot(x=class_counts.index, y=class_counts.values, palette='viridis', hue=class_counts.index, legend=False)
7
8 plt.xlabel('Classes', fontsize=12)
9 plt.ylabel('Count', fontsize=12)
10 plt.title('Class Distribution', fontsize=16)
11
12 for i, count in enumerate(class_counts.values):
13     plt.text(i, count, str(count), ha='center', va='bottom', fontsize=10)
14
15 plt.show()
✓ [1502] 75ms

```



### 3. Data pre-processing:

- **Null values:**

There are null values in our chosen dataset.

```

1 data.isnull().sum()
[142]

Name                                0
Team Name                           0
Position                             0
Age                                  0
Season                              0
FIFA rating                          0
Injury                              0
Date of Injury                       0
Date of return                       0
Match1_before_injury_Result          65
Match1_before_injury_Opposition       65
Match1_before_injury_GD               65
Match1_before_injury_Player_rating    67
Match2_before_injury_Result          101
Match2_before_injury_Opposition       101
Match2_before_injury_GD               101
Match2_before_injury_Player_rating    102
Match3_before_injury_Result          157
Match3_before_injury_Opposition       157
Match3_before_injury_GD               157

```

To remedy this we dropped the columns that have too many null values along with the features that we do not want to work with in our model training.

```

1 print("shape of dataframe before dropping:", data.shape)
2
3 data = data.drop(['Name', 'Match2_after_injury_Result', 'Match2_after_injury_Opposition', 'Match2_after_injury_GD',
4     'Match2_after_injury_Player_rating', 'Match3_after_injury_Result', 'Match3_after_injury_Opposition',
5     'Match3_after_injury_GD', 'Match3_after_injury_Player_rating', 'Match3_missed_match_Result',
6     'Match3_missed_match_Opposition', 'Match3_missed_match_GD', 'Match1_after_injury_GD',
7     'Match1_after_injury_Player_rating'], axis=1)
8 print("shape of dataframe after dropping:", data.shape)
[143]

shape of dataframe before dropping: (656, 42)
shape of dataframe after dropping: (656, 28)

```



We also dropped the rows where the target value which we are trying to predict is null.

```
1 print("shape of dataframe before dropping:", data.shape)
2
3 target_column = 'Match1_after_injury_Result'
4
5 data = data.dropna(subset=[target_column])
6
7 print("shape of dataframe after dropping:", data.shape)
✓ [11] < 10 ms

shape of dataframe before dropping: (656, 28)
shape of dataframe after dropping: (515, 28)
```

- **Categorical values:**

Machine Learning models can only interpret numerical or qualitative data. However, our dataset has many features that are categorical.

```
1 data.info()
```

```
✓ [5] < 10 ms
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 656 entries, 0 to 655
```

```
Data columns (total 42 columns):
```

#	Column	Non-Null Count	Dtype
0	Name	656 non-null	object
1	Team Name	656 non-null	object
2	Position	656 non-null	object
3	Age	656 non-null	int64
4	Season	656 non-null	object
5	FIFA rating	656 non-null	int64
6	Injury	656 non-null	object
7	Date of Injury	656 non-null	object
8	Date of return	656 non-null	object
9	Match1_before_injury_Result	591 non-null	object
10	Match1_before_injury_Opposition	591 non-null	object
11	Match1_before_injury_GD	591 non-null	float64
12	Match1_before_injury_Player_rating	589 non-null	object
13	Match2_before_injury_Result	555 non-null	object
14	Match2_before_injury_Opposition	555 non-null	object

For the models to be able to understand these categorical features, we need to encode these data to numeric values.

We used one hot encoding for those features where each classification value is of the same importance.

One Hot Encoding

```
1 data = pd.get_dummies(data, columns=['Team Name', 'Position', 'Season', 'Injury', 'Match1_before_injury_Opposition',  
2                                     'Match2_before_injury_Opposition', 'Match3_before_injury_Opposition',  
3                                     'Match1_after_injury_Opposition', 'Match1_missed_match_Opposition',  
4                                     'Match2_missed_match_Opposition'], drop_first=True)  
5 print(data.shape)    data  
6 data.head()
```

```
✓ [56] 12ms
```

```
(515, 311)
```

Team Name_Aston Villa	Team Name_Brentford	Team Name_Burnley	Team Name_Everton	Team Name_Man United	Team Name_Newcastle	Team Name_Tottenham	Position_Center Back
False	False	False	False	False	• True	False	• True
False	False	False	False	False	• True	False	• True
False	False	False	False	False	• True	False	False
False	False	False	False	False	• True	False	False

We used label encoding for the features where the categorical values hold some significance. For example winning is better than a draw or a loss.

Label Encoding

```

1 data['Match1_before_injury_Result'].unique()
2 ✓ [57] < 10 ms
3 array(['draw', 'lose', 'win', nan], dtype=object)
4
5 mappings = {'Match1_before_injury_Result': {'draw': 1, 'lose': 0, 'win': 2},
6           'Match2_before_injury_Result': {'draw': 1, 'lose': 0, 'win': 2},
7           'Match3_before_injury_Result': {'draw': 1, 'lose': 0, 'win': 2},
8           'Match1_after_injury_Result': {'draw': 1, 'lose': 0, 'win': 2},
9           'Match1_missed_match_Result': {'draw': 1, 'lose': 0, 'win': 2},
10          'Match2_missed_match_Result': {'draw': 1, 'lose': 0, 'win': 2}}
11
12 for column, mapping in mappings.items():
13     data[column] = data[column].map(mapping)
14
15 data.head()
16 ✓ [58] 14ms

```

	Age	FIFA rating	Date of Injury	Date of return	Match1_before_injury_Result	Match1_before_injury_GO	Match1_before_injury_Player_rating	Match2_before_injury_Result
0	26	77	Nov 9, 2019	Jan 13, 2020	1.0	0.0	7.4	2.0
2	28	79	Jan 2, 2020	Jan 17, 2020	0.0	-3.0	4.9	0.0
5	30	75	Dec 9, 2019	Jan 10, 2020	1.0	0.0	6	2.0
7	26	77	Sep 22, 2019	Oct 5, 2019	1.0	0.0	6.7	0.0
8	26	77	Dec 15, 2019	Dec 25, 2019	2.0	2.0	6	2.0

- **Date & Time:**

The features indicating a date in our datasets are stored as strings. For the machine learning model to understand the data we need to convert that to datetime type data and then to a numeric value. For our dataset we felt that the month holds the most important information. So, we just kept the month value as a number.

```

1 datetime_columns = ['Date of Injury', 'Date of return']
2
3 for col in datetime_columns:
4     data[col] = pd.to_datetime(data[col], errors='coerce')
5     data[col] = data[col].apply(lambda row: row.month)
6
7 data.head()
8 ✓ [60] 14ms

```

	Age	FIFA rating	Date of Injury	Date of return
0	26	77	11.0	1.0
2	28	79	1.0	1.0
5	30	75	12.0	1.0
7	26	77	9.0	10.0
8	26	77	12.0	12.0

- **Objects to numeric:**

There are some features left which are in object data type. We converted those to numeric values as well.

```
data['Match1_before_injury_Player_rating'] = pd.to_numeric(data['Match1_before_injury_Player_rating'], errors='coerce')
data['Match2_before_injury_Player_rating'] = pd.to_numeric(data['Match2_before_injury_Player_rating'], errors='coerce')
data['Match3_before_injury_Player_rating'] = pd.to_numeric(data['Match3_before_injury_Player_rating'], errors='coerce')
✓ [59] < 10 ms
```

## 4. Feature scaling:

We believe that our particular dataset does not require any scaling because there are no such features in our dataset that have overwhelming variance compared to other features.

We experimented with MinMaxScaler and Standard Scaler and found that our models accuracy decreases if we use these scalers.

## 5. Dataset splitting:

We splitted our dataset into a Train set and Test set where the train set contains 70% of the original dataset. Stratified splitting was used to ensure that each unique categorical target data was adequately present in our train and test set.

```

1 from sklearn.model_selection import train_test_split
2
3 target_column = 'Match1_after_injury_Result'
4
5 x = data.drop(target_column, axis=1).fillna(0)
6 y = data[target_column]
7
8 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=42, stratify=y)
✓ [64] < 10 ms

```

```

1 x_train.info()
✓ [65] < 10 ms

<class 'pandas.core.frame.DataFrame'>
Index: 360 entries, 44 to 617
Columns: 310 entries, Age to Match2_missed_match_Opposition_Wolves
dtypes: bool(293), float64(15), int64(2)
memory usage: 153.6 KB

```

```

1 x_test.info()
✓ [66] < 10 ms

<class 'pandas.core.frame.DataFrame'>
Index: 155 entries, 559 to 538
Columns: 310 entries, Age to Match2_missed_match_Opposition_Wolves
dtypes: bool(293), float64(15), int64(2)
memory usage: 66.1 KB

```

```

1 y_train.info()
✓ [67] < 10 ms

<class 'pandas.core.series.Series'>
Index: 360 entries, 44 to 617
Series name: Match1_after_injury_Result
Non-Null Count  Dtype
-----
360 non-null    int64
dtypes: int64(1)
memory usage: 5.6 KB

```

```

1 y_test.info()
✓ [68] < 10 ms

<class 'pandas.core.series.Series'>
Index: 155 entries, 559 to 538
Series name: Match1_after_injury_Result
Non-Null Count  Dtype
-----
155 non-null    int64
dtypes: int64(1)
memory usage: 2.4 KB

```

## 6. Model training & testing:

- **Decision Tree:**

We trained a decision tree model using our train set and tested with the test set and found out that this particular model has 42% accuracy for our dataset.

```
1 from sklearn.tree import DecisionTreeClassifier
2
3 model_decision_tree = DecisionTreeClassifier(random_state=42)
4 model_decision_tree.fit(x_train, y_train)
5
6 y_pred_Ddecision_Tree = model_decision_tree.predict(x_test)
7
8 accuracy_Ddecision_Tree = model_decision_tree.score(x_test, y_test)
9 accuracy_Ddecision_Tree = round(accuracy_Ddecision_Tree, 2)
10 print(accuracy_Ddecision_Tree)
✓ [69] 15ms
0.42
```

- **Random Forest:**

We trained a random forest model using our train set and tested with the test set and found out that this particular model has 50% accuracy for our dataset.

```
1 from sklearn.ensemble import RandomForestClassifier
2
3 model_random_forest = RandomForestClassifier(random_state=42)
4 model_random_forest.fit(x_train, y_train)
5
6 y_pred_Random_Forest = model_random_forest.predict(x_test)
7
8 accuracy_Random_Forest = model_random_forest.score(x_test, y_test)
9 accuracy_Random_Forest = round(accuracy_Random_Forest, 2)
10 print(accuracy_Random_Forest)
✓ [70] 112ms
0.5
```

- **Hist Gradient Boosting Classifier:**

We trained a hist gradient boosting classifier model using our train set and tested with the test set and found out that this particular model has 52% accuracy for our dataset.

```
1 from sklearn.ensemble import HistGradientBoostingClassifier
2
3 model_HistGradientBoostingClassifier = HistGradientBoostingClassifier(random_state=42)
4 model_HistGradientBoostingClassifier.fit(x_train, y_train)
5
6 y_pred_HistGradientBoostingClassifier = model_HistGradientBoostingClassifier.predict(x_test)
7
8 accuracy_HistGradientBoostingClassifier = model_HistGradientBoostingClassifier.score(x_test, y_test)
9 accuracy_HistGradientBoostingClassifier = round(accuracy_HistGradientBoostingClassifier, 2)
10 print(accuracy_HistGradientBoostingClassifier)
✓ [71] 1s 399ms
0.52
```

- **KNN:**

We trained a K nearest neighbor model using our train set and tested with the test set and found out that this particular model has 39% accuracy for our dataset.

```
1 from sklearn.neighbors import KNeighborsClassifier
2
3 modelKNeighborsClassifier = KNeighborsClassifier(n_neighbors=8)
4 modelKNeighborsClassifier.fit(x_train, y_train)
5
6 y_pred_KNeighborsClassifier = modelKNeighborsClassifier.predict(x_test)
7 accuracy_KNeighborsClassifier = modelKNeighborsClassifier.score(x_test, y_test)
8 accuracy_KNeighborsClassifier = round(accuracy_KNeighborsClassifier, 2)
9 print(accuracy_KNeighborsClassifier)
✓ [72] 31ms
0.39
```

- **Neural Networks:**

Finally, we implemented a neural network model using tensorflow and it provided an accuracy of 47%.

```

1 num_classes = len(np.unique(y))
2
3 if num_classes > 2:
4     y_train_NN = tf.keras.utils.to_categorical(y_train, num_classes)
5     y_test_NN = tf.keras.utils.to_categorical(y_test, num_classes)
6
7 model = Sequential([
8     Input(shape=(x_train.shape[1],)),
9     Dense(64, activation='relu'),
10    Dense(32, activation='relu'),
11    Dense(num_classes, activation='softmax' if num_classes > 2 else 'sigmoid')
12 ])
✓ [1524] 13ms

```

```

1 model.compile(
2     optimizer='adam',
3     loss='categorical_crossentropy' if num_classes > 2 else 'binary_crossentropy',
4     metrics=['accuracy'])
5 )
✓ [1525] < 10 ms

```

```

1 history = model.fit(
2     x_train,
3     y_train_NN,
4     epochs=50,
5     batch_size=32,
6     validation_split=0.3
7 )
✓ [1526] 3s 314ms

```

```

Epoch 41/50
8/8 ————— 0s 5ms/step - accuracy: 0.9792 - loss: 0.2494 - val_accuracy: 0.3853 - val_loss: 1.3981
Epoch 42/50
8/8 ————— 0s 5ms/step - accuracy: 0.9914 - loss: 0.2427 - val_accuracy: 0.3670 - val_loss: 1.4744
Epoch 43/50
8/8 ————— 0s 6ms/step - accuracy: 0.9517 - loss: 0.2707 - val_accuracy: 0.4220 - val_loss: 1.4440
Epoch 44/50
8/8 ————— 0s 6ms/step - accuracy: 0.9684 - loss: 0.2318 - val_accuracy: 0.4037 - val_loss: 1.5071
Epoch 45/50
8/8 ————— 0s 6ms/step - accuracy: 0.9801 - loss: 0.2236 - val_accuracy: 0.3394 - val_loss: 1.4786
Epoch 46/50
8/8 ————— 0s 5ms/step - accuracy: 0.9844 - loss: 0.1927 - val_accuracy: 0.4220 - val_loss: 1.4996
Epoch 47/50
8/8 ————— 0s 6ms/step - accuracy: 0.9927 - loss: 0.1848 - val_accuracy: 0.3945 - val_loss: 1.5122
Epoch 48/50
8/8 ————— 0s 5ms/step - accuracy: 0.9840 - loss: 0.1924 - val_accuracy: 0.3945 - val_loss: 1.5101
Epoch 49/50
8/8 ————— 0s 6ms/step - accuracy: 0.9904 - loss: 0.1804 - val_accuracy: 0.4037 - val_loss: 1.5080
Epoch 50/50
8/8 ————— 0s 6ms/step - accuracy: 0.9916 - loss: 0.1722 - val_accuracy: 0.4312 - val_loss: 1.5838

```

```

1 test_loss, test_accuracy_NN = model.evaluate(x_test, y_test_NN)
2 test_accuracy_NN = round(test_accuracy_NN, 2)
3 print(f"Test Accuracy: {test_accuracy_NN}")
✓ [1527] 50ms

```

```

5/5 ————— 0s 3ms/step - accuracy: 0.4365 - loss: 1.4133
Test Accuracy: 0.47

```

```

1 y_pred_NN = model.predict(x_test)
✓ [1528] 69ms

```


```

5/5 ————— 0s 6ms/step

```




The model loss and accuracy graph for this neural network model was:



## 7. Model selection/Comparison analysis:

- Bar chart showcasing prediction accuracy of all models:

```
1 models = ['Decision\nTree', 'Random\nForest', 'KNN', 'Hist\nGradient\nBoosting', 'Neural\nNetwork']
2 accuracies = [accuracy_Decision_Tree, accuracy_Random_Forest, accuracy_KNeighborsClassifier,
3               accuracy_HistGradientBoostingClassifier, test_accuracy_NN]
4
5 plt.bar(models, accuracies, tick_label=models)
6 plt.title('Model Accuracy')
7 plt.xlabel('Model')
8 plt.ylabel('Accuracy')
9 plt.show()
✓ [1530] 62ms
```



- Precision, recall comparison of each model:

(a) Decision Tree:

```
report_decision_tree = classification_report(y_test, y_pred_Decision_Tree)
print(report_decision_tree)
```

✓ [432] < 10 ms

	precision	recall	f1-score	support
0	0.39	0.45	0.42	53
1	0.35	0.29	0.32	38
2	0.48	0.47	0.48	64
accuracy			0.42	155
macro avg	0.41	0.40	0.40	155
weighted avg	0.42	0.42	0.42	155

### (b) Random Forest:

```
1 report_random_forest = classification_report(y_test, y_pred_Random_Forest)
2 print(report_random_forest)
```

✓ [433] < 10 ms

	precision	recall	f1-score	support
0	0.49	0.49	0.49	53
1	0.67	0.16	0.26	38
2	0.49	0.72	0.59	64
accuracy			0.50	155
macro avg	0.55	0.46	0.44	155
weighted avg	0.54	0.50	0.47	155

### (c) KNN:

```
1 report_KNeighborsClassifier = classification_report(y_test, y_pred_KNeighborsClassifier)
2 print(report_KNeighborsClassifier)
```

✓ [434] < 10 ms

	precision	recall	f1-score	support
0	0.38	0.58	0.46	53
1	0.35	0.18	0.24	38
2	0.43	0.36	0.39	64
accuracy			0.39	155
macro avg	0.39	0.38	0.36	155
weighted avg	0.39	0.39	0.38	155

#### (d) Hist Gradient Boosting:

```
1 report_HistGradientBoostingClassifier = classification_report(y_test, y_pred_HistGradientBoostingClassifier)
2 print(report_HistGradientBoostingClassifier)
✓ [435] < 10 ms
```

	precision	recall	f1-score	support
0	0.48	0.51	0.50	53
1	0.48	0.39	0.43	38
2	0.56	0.59	0.58	64
accuracy			0.52	155
macro avg	0.51	0.50	0.50	155
weighted avg	0.51	0.52	0.51	155

#### (e) Neural Network:

```
1 y_pred_prob = model.predict(x_test)
2 if num_classes > 2:
3     y_pred_class = np.argmax(y_pred_prob, axis=1)
4     y_test_class = np.argmax(y_test_NN, axis=1)
5 else:
6     y_pred_class = (y_pred_prob > 0.5).astype(int).flatten()
7     y_test_class = y_test.flatten()
8
9 report_NN = classification_report(y_test_class, y_pred_class)
10 print(report)
11
✓ [1535] 51ms
```

1/5 ————— 0s 9ms/step5/5 ————— 0s 3m

	precision	recall	f1-score	support
0	0.42	0.26	0.33	53
1	0.32	0.26	0.29	38
2	0.49	0.70	0.58	64
accuracy			0.45	155
macro avg	0.41	0.41	0.40	155
weighted avg	0.43	0.45	0.42	155


- **Confusion matrix:**

#### (a) Decision Tree:

```

1  conf_matrix_decision_tree = confusion_matrix(y_test, y_pred_Decision_Tree)
2
3  plt.figure(figsize=(4, 3))
4  sns.heatmap(conf_matrix_decision_tree, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y_test),
5              yticklabels=np.unique(y_test))
6  plt.xlabel('Predicted Label')
7  plt.ylabel('True Label')
8  plt.title('Confusion Matrix')
9  plt.show()
✓ [1536] 75ms

```




## (b) Random Forest:

```


1  conf_matrix_random_forest = confusion_matrix(y_test, y_pred_Random_Forest)
2
3  plt.figure(figsize=(4, 3))
4  sns.heatmap(conf_matrix_random_forest, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y_test),
5              yticklabels=np.unique(y_test))
6  plt.xlabel('Predicted Label')
7  plt.ylabel('True Label')
8  plt.title('Confusion Matrix')
9  plt.show()
✓ [1537] 71ms

```




### (c) KNN:

```
1 conf_matrix_KNeighborsClassifier = confusion_matrix(y_test, y_pred_KNeighborsClassifier)
2
3 plt.figure(figsize=(4, 3))
4 sns.heatmap(conf_matrix_KNeighborsClassifier, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y_test),
5             yticklabels=np.unique(y_test))
6 plt.xlabel('Predicted Label')
7 plt.ylabel('True Label')
8 plt.title('Confusion Matrix')
9 plt.show()
✓ [1538] 75ms
```




### (d) Hist Gradient Boosting:

```
1 conf_matrix_hist_gradient_boosting = confusion_matrix(y_test, y_pred_HistGradientBoostingClassifier)
2
3 plt.figure(figsize=(4, 3))
4 sns.heatmap(conf_matrix_hist_gradient_boosting, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y_test),
5             yticklabels=np.unique(y_test))
6 plt.xlabel('Predicted Label')
7 plt.ylabel('True Label')
8 plt.title('Confusion Matrix')
9 plt.show()
✓ [1539] 75ms
```



### (e) Neural Network:

```
1 conf_matrix_NN = confusion_matrix(y_test_class, y_pred_class)
2
3 plt.figure(figsize=(4, 3))
4 sns.heatmap(conf_matrix_NN, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y_test),
5             yticklabels=np.unique(y_test))
6 plt.xlabel('Predicted Label')
7 plt.ylabel('True Label')
8 plt.title('Confusion Matrix')
9 plt.show()
✓ [1540] 73ms
```



## 8. Conclusion:

From the performance metrics of the above models we can conclude that for our particular dataset the Hist Gradient Boosting model was the most successful among the used models. The accuracy of prediction for all the models were not satisfactory. The reason behind this can be the lack of correlation between the features and the target variable which is to be expected as predicting football match results is not an easy task. The result of a football match depends on the performance of each player of a team and the tactics of the coaching staff on the pitch. Predicting a match result with a finite number of attributes with great accuracy is almost an impossible task. This is the thing that makes a football match interesting and enjoyable.